

libmemkind探究（一）——让jemalloc管理指定的空间

最终我自己写的分配器——在《[面向非易失存储器（NVM）的内存分配器——libnvmalloc](#)》中——被弃用了，以安全为由。于是必须得使用libmemkind，这是Intel官方的一个内存分配器函数库，为pmem文件系统设计的。从一开始，我就知道这个库是基于jemalloc的。但是一直不知道具体它是如何让jemalloc乖乖地在指定的内存空间上管理的。

=====阶段一：jemalloc提供的接口
=====

于是我就随便下载了[jemalloc-4.0.3](#)的源码包，先去看了doc/jemalloc.html。在阅读jemalloc的doc/jemalloc.html时，有一段引起了我的兴趣：

```
"arena.<i>.chunk_hooks" (chunk_hooks_t) rw

Get or set the chunk management hook functions for arena <i>. The functions must be capable of operating on all extant chunks associated with arena <i>, usually by passing unknown chunks to the replaced functions. In practice, it is feasible to control allocation for arenas created via "arenas.extend" such that all chunks originate from an application-supplied chunk allocator (by setting custom chunk hook functions just after arena creation), but the automatically created arenas may have already created chunks prior to the application having an opportunity to take over chunk allocation.

typedef struct {
    chunk_alloc_t      *alloc;
    chunk_dalloc_t     *dalloc;
    chunk_commit_t     *commit;
    chunk_decommit_t   *decommit;
    chunk_purge_t      *purge;
    chunk_split_t      *split;
    chunk_merge_t      *merge;
} chunk_hooks_t;
```

也就是说，我可以向jemalloc注册一组函数，来控制chunk的申请/释放行为。我知道，一个arena就像一个命名空间，不同的arena管理不同的内存区域。但是arena毕竟只是一个虚的概念，真正承载用户数据的是chunk。现在能够指定chunk怎么申请、怎么释放，那么我就能控制chunk的内存空间的来源。文中也说了，已经自动创建的arena可能已经先于用户注册自定义函数而申请了一些chunk，那么这些chunk肯定是从DDR里面分来的。

为了确保我所使用的内存空间来自我指定的地方，那么我就需要新建arena，然后给该arena绑定自定义的chunk hooks，来自定义chunk的行为。文中也提到，arena可以通过“arenas.extend”创建。于是我跳到那一段：

```
"arenas.extend" (unsigned) r-

Extend the array of arenas by appending a new arena, and returning the new arena index.
```

那么在用户通过诸如je_malloc()函数申请内存时，又如何确保内存来自这个新建的arena呢？我有在文档中搜寻，果然找到了：

Non-standard API

The `mallocx()`, `rallocx()`, `xallocx()`, `sallocx()`, `dallocx()`, `sdallocx()`, and `nallocx()` functions all have a `flags` argument that can be used to specify options. The functions only check the options that are contextually relevant. Use bitwise or (`|`) operations to specify one or more of the following:

MALLOCX_LG_ALIGN(*la*)

Align the memory allocation to start at an address that is a multiple of $(1 \ll la)$. This macro does not validate that *la* is within the valid range.

MALLOCX_ALIGN(*a*)

Align the memory allocation to start at an address that is a multiple of *a*, where *a* is a power of two. This macro does not validate that *a* is a power of 2.

MALLOCX_ZERO

Initialize newly allocated memory to contain zero bytes. In the growing reallocation case, the real size prior to reallocation defines the boundary between untouched bytes and those that are initialized to contain zero bytes. If this macro is absent, newly allocated memory is uninitialized.

MALLOCX_TCACHE(*tc*)

Use the thread-specific cache (tcache) specified by the identifier *tc*, which must have been acquired via the ["tcache.create"](#) mallctl. This macro does not validate that *tc* specifies a valid identifier.

MALLOCX_TCACHE_NONE

Do not use a thread-specific cache (tcache). Unless `MALLOCX_TCACHE(tc)` or `MALLOCX_TCACHE_NONE` is specified, an automatically managed tcache will be used under many circumstances. This macro cannot be used in the same `flags` argument as `MALLOCX_TCACHE(tc)`.

MALLOCX_ARENA(*a*)

Use the arena specified by the index *a*. This macro has no effect for regions that were allocated via an arena other than the one specified. This macro does not validate that *a* specifies an arena index in the valid range.

因此大致思路有了：

- 在初始化时，使用`je_mallctl("arenas.extend",...)`新建一个arena，并且得到新的arena的编号`arena_index`；
- 填充一个`chunk_hooks_t`结构体，并且使用`je_mallctl("arenas.<arena_index>.chunk_hooks",...)`为这个arena注册自定义chunk行为；
- 在之后的分配中，使用`je_mallocx(size, MALLOCX_ARENA(arena_index))`来要求jemalloc从该arena中分配空间。

=====阶段二：libmemkind利用该接口

=====

原来jemalloc都留好了现成的接口了，我猜libmemkind就是这么实现的吧？为了验证猜想，我下载了libmemkind v1.5.0，并开始查看其源码。果然，我发现了这么一条调用链：

```
memkind_pmem_create(kind, ops, name);           // 用户调用它，
memkind_arena_create_map(kind, &pmem_chunk_hooks); // 使用pmem_ch
// something
unsigned i = 0;
for(i=0; i < kind->arena_map_len; i++) {
    //create new arena with consecutive index
    unsigned arena_index;
    err = jemk_mallctl("arenas.extend", (void*)&arena_index, &unsi
    if(err) {
        goto exit;
    }
    //store index of first arena
    if(i == 0) {
        kind->arena_zero = arena_index;
    }
    //setup chunk_hooks for newly created arena
    char cmd[64];
    snprintf(cmd, sizeof(cmd), "arena.%u.chunk_hooks", arena_index
```

```

    err = jemk_mallocctl(cmd, NULL, NULL, (void*)hooks, chunk_hooks_
    if(err) {
        goto exit;
    }
    arena_registry_g[arena_index] = kind;
}

```

也就是说，新建了多个arena，并且为他们注册了pmem_chunk_hooks。而这个pmem_chunk_hooks就是一个chunk_hooks_t结构体，定义了怎么从pmem文件申请/释放chunk：

```

void *pmem_chunk_alloc(void *chunk, size_t size, size_t alignment,
    bool *zero, bool *commit, unsigned arena_ind)
{
    int err;
    void *addr = NULL;

    if (chunk != NULL) {
        /* not supported */
        goto exit;
    }

    struct memkind *kind;
    kind = get_kind_by_arena(arena_ind);
    if (kind == NULL) {
        return NULL;
    }

    err = memkind_check_available(kind);
    if (err) {
        goto exit;
    }

    addr = memkind_pmem_mmap(kind, chunk, size);

    if (addr != MAP_FAILED) {
        *zero = true;
        *commit = true;

        /* XXX - check alignment */
    } else {
        addr = NULL;
    }

    exit:
    return addr;
}

bool pmem_chunk_dalloc(void *chunk, size_t size, bool committed,
    unsigned arena_ind)
{

```

```
/* do nothing - report failure (opt-out) */
return true;
}

bool pmem_chunk_commit(void *chunk, size_t size, size_t offset, size_t
    unsigned arena_ind)
{
    /* do nothing - report success */
    return false;
}

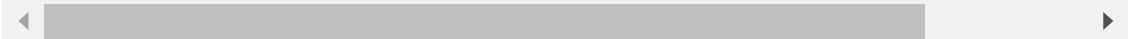
bool pmem_chunk_decommit(void *chunk, size_t size, size_t offset, size_t
    unsigned arena_ind)
{
    /* do nothing - report failure (opt-out) */
    return true;
}

bool pmem_chunk_purge(void *chunk, size_t size, size_t offset, size_t
    unsigned arena_ind)
{
    /* do nothing - report failure (opt-out) */
    return true;
}

bool pmem_chunk_split(void *chunk, size_t size, size_t size_a, size_t
    bool committed, unsigned arena_ind)
{
    /* do nothing - report success */
    return false;
}

bool pmem_chunk_merge(void *chunk_a, size_t size_a, void *chunk_b,
    size_t size_b, bool committed, unsigned arena_ind)
{
    /* do nothing - report success */
    return false;
}

static chunk_hooks_t pmem_chunk_hooks = {
    pmem_chunk_alloc,
    pmem_chunk_dalloc,
    pmem_chunk_commit,
    pmem_chunk_decommit,
    pmem_chunk_purge,
    pmem_chunk_split,
    pmem_chunk_merge
};
```



可以看到，在`pmem_chunk_alloc`时，使用了`memkind_pmem_mmap()`来分配空间。那么再看这个`memkind_pmem_mmap()`函数：

```
MEMKIND_EXPORT void *memkind_pmem_mmap(struct memkind *kind, void *addr,
{
    struct memkind_pmem *priv = kind->priv;
    void *result;

    if (pthread_mutex_lock(&priv->pmem_lock) != 0)
        assert(0 && "failed to acquire mutex");

    if (priv->offset + size > priv->max_size) {
        pthread_mutex_unlock(&priv->pmem_lock);
        return MAP_FAILED;
    }

    if ((errno = posix_fallocate(priv->fd, priv->offset, size)) != 0)
        pthread_mutex_unlock(&priv->pmem_lock);
        return MAP_FAILED;
    }

    result = priv->addr + priv->offset;
    priv->offset += size;

    pthread_mutex_unlock(&priv->pmem_lock);

    return result;
}
```

从它的代码逻辑可以看出：该函数会试图从`pmem`文件中分配`[offset, offset + size)`这么一段空间。为了确保空间能有物理空间对应，使用了`posix_fallocate()`对文件“打洞”。成功后，就把该段空间分配出去，并且`offset`增加`size`。可以看到，`libmemkind`只能不断向后分配空间，但是却不能释放空间（当然，这不是什么大问题）。

=====阶段三：自己写程序验证
=====

`libmemkind`毕竟还有很多其他的内容，并不能很好的凸显`jemalloc`提供的接口的使用方式。为了验证我的想法确实可行，还是自己写程序最靠谱。我写了一个只有140行的测试程序，但是可以很好地显示出，`jemalloc`确实在管理我模拟的一段空间。

test.c

```
#include <stdio.h>
#include <jemalloc.h>
#include <sys/mman.h>

#define PAGE_SIZE 4096

// 待管理的空间大小
size_t mem_size = 1 << 30;
// 待管理的空间的首地址
```

```

void* base;
// 已经分配出去的大小
size_t allocated = 0;

//+++++HOOK BEGIN+++++

// 从这里开始定义jemalloc管理chunk的hook函数，自定义chunk管理行为
// 可以参考jemalloc/doc/jemalloc.html中arena.<i>.chunk_hooks一段

// 当jemalloc发现chunk不够用了，会callback此函数索要空间
// chunk大小在编译时配置（原版jemalloc-4.0.3默认人2M，libmemkind配置为4M）
void* my_chunk_alloc(void *chunk, size_t size, size_t alignment, bool
{
    printf("my_chunk_alloc(chunk = %p, size = %lu, alignment = %lu)\n",
        if(size % alignment)
            return NULL;
        if(allocated + size > mem_size)
            return NULL;
        if(chunk && chunk != base + allocated)
            return NULL;
        void* addr = base + allocated;
        allocated += size;
        return addr;
    }

// 返回true表示该内存可以继续使用
bool my_chunk_dalloc(void *chunk, size_t size, bool committed, unsigned
{
    return true;
}

// 返回false表示内存充足
bool my_chunk_commit(void *chunk, size_t size, size_t offset, size_t l
{
    return false;
}

// 返回true表示该内存即使释放了，也是与物理内存对应的，可以重用
bool my_chunk_decommit(void *chunk, size_t size, size_t offset, size_t
{
    return true;
}

// 返回true表示该段地址空间被重用后不会清空
bool my_chunk_purge(void *chunk, size_t size, size_t offset, size_t le
{
    return true;
}

bool my_chunk_split(void *chunk, size_t size, size_t size_a, size_t si
{
    return false;
}

```

```

bool my_chunk_merge(void *chunk_a, size_t size_a, void *chunk_b, size_t
{
    return false;
}

chunk_hooks_t my_chunk_hooks =
{
    .alloc = my_chunk_alloc,
    .dalloc = my_chunk_dalloc,
    .commit = my_chunk_commit,
    .decommit = my_chunk_decommit,
    .purge = my_chunk_purge,
    .split = my_chunk_split,
    .merge = my_chunk_merge,
};
//-----HOOK END-----

int main()
{
    // 模拟一段NVM空间, 或者任何一段用户待管理的空间 (这里有1GB)
    base = mmap(0, mem_size, PROT_READ | PROT_WRITE, MAP_ANONYMOUS | MAP_PRIVATE,
    if(base == MAP_FAILED)
    {
        printf("mmap() failed!\n");
        return 1;
    }
    printf("base = %p\n", base);

    // 用je_mallctl('arenas.extend')命令创建一个arean,
    // 参考jemalloc/doc/jemalloc.html中arenas.extend一段
    unsigned arena_index;
    size_t unsigned_size = sizeof(unsigned int);
    if(je_mallctl("arenas.extend", (void*)&arena_index, &unsigned_size, 0, 0) != 0)
    {
        printf("je_mallctl('arenas.extend') failed!\n");
        return 1;
    }
    printf("arena_index = %u\n", arena_index);

    //为这个arena绑定我们自定义的chunk hook, 于是该arena就会按我们的方式去管理内存了
    // 参考jemalloc/doc/jemalloc.html中arena.<i>.chunk_hooks一段
    char cmd[64];
    sprintf(cmd, "arena.%u.chunk_hooks", arena_index);
    if(je_mallctl(cmd, NULL, NULL, (void*)&my_chunk_hooks, sizeof(chunk_hooks_t), 0) != 0)
    {
        printf("je_mallctl('%s') failed!\n", cmd);
        return 1;
    }

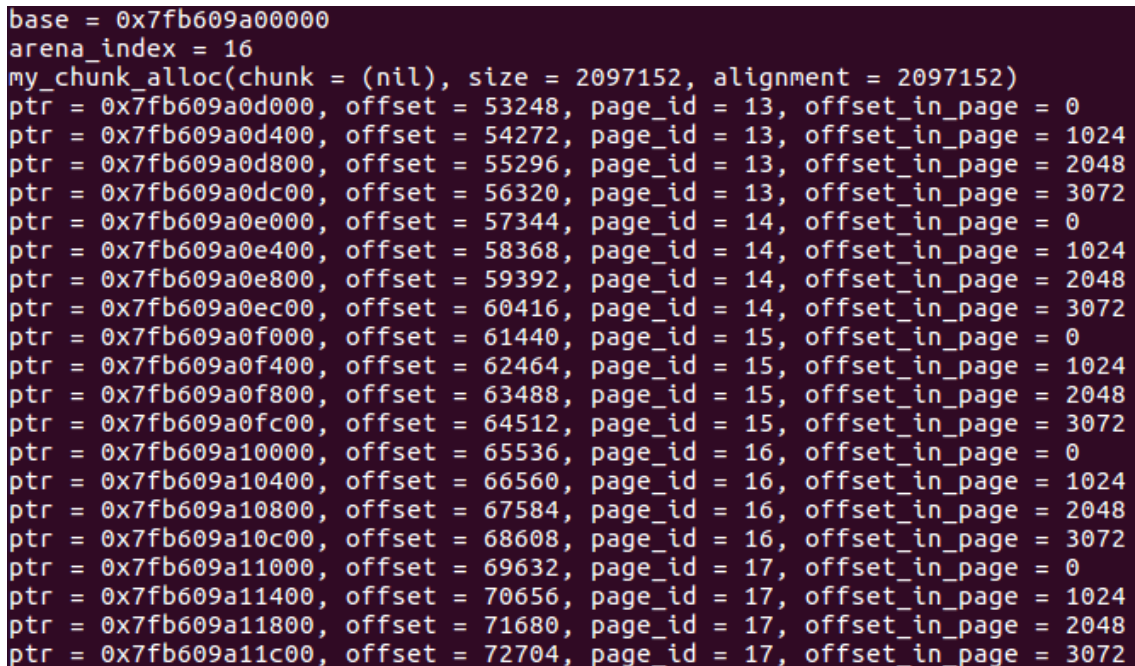
    // 接着试试分配内存, 使用je_mallocx()、je_callocx()等以x结尾的函数, 指定内存大小
    // 那么就会在我们创建的arena中分配内存
    for(size_t i = 0; i < 1024; i++)

```

```
{
    void* ptr = je_mallocx(1024, MALLOCX_ARENA(arena_index));
    // 算一算偏移量
    size_t offset = ptr - base;
    // 是待管理空间的第几个page
    size_t page_id = offset / 4096;
    // page内偏移多少
    size_t offset_in_page = offset % 4096;
    printf("ptr = %p, offset = %lu, page_id = %lu, offset_in_page\n", ptr, offset, page_id, offset_in_page);
}

if(munmap(base, mem_size) != 0)
{
    printf("munmap() failed!\n");
    return 1;
}
return 0;
}
```

运行结果如下:



```
base = 0x7fb609a00000
arena_index = 16
my_chunk_alloc(chunk = (nil), size = 2097152, alignment = 2097152)
ptr = 0x7fb609a0d000, offset = 53248, page_id = 13, offset_in_page = 0
ptr = 0x7fb609a0d400, offset = 54272, page_id = 13, offset_in_page = 1024
ptr = 0x7fb609a0d800, offset = 55296, page_id = 13, offset_in_page = 2048
ptr = 0x7fb609a0dc00, offset = 56320, page_id = 13, offset_in_page = 3072
ptr = 0x7fb609a0e000, offset = 57344, page_id = 14, offset_in_page = 0
ptr = 0x7fb609a0e400, offset = 58368, page_id = 14, offset_in_page = 1024
ptr = 0x7fb609a0e800, offset = 59392, page_id = 14, offset_in_page = 2048
ptr = 0x7fb609a0ec00, offset = 60416, page_id = 14, offset_in_page = 3072
ptr = 0x7fb609a0f000, offset = 61440, page_id = 15, offset_in_page = 0
ptr = 0x7fb609a0f400, offset = 62464, page_id = 15, offset_in_page = 1024
ptr = 0x7fb609a0f800, offset = 63488, page_id = 15, offset_in_page = 2048
ptr = 0x7fb609a0fc00, offset = 64512, page_id = 15, offset_in_page = 3072
ptr = 0x7fb609a10000, offset = 65536, page_id = 16, offset_in_page = 0
ptr = 0x7fb609a10400, offset = 66560, page_id = 16, offset_in_page = 1024
ptr = 0x7fb609a10800, offset = 67584, page_id = 16, offset_in_page = 2048
ptr = 0x7fb609a10c00, offset = 68608, page_id = 16, offset_in_page = 3072
ptr = 0x7fb609a11000, offset = 69632, page_id = 17, offset_in_page = 0
ptr = 0x7fb609a11400, offset = 70656, page_id = 17, offset_in_page = 1024
ptr = 0x7fb609a11800, offset = 71680, page_id = 17, offset_in_page = 2048
ptr = 0x7fb609a11c00, offset = 72704, page_id = 17, offset_in_page = 3072
```

可以看到my_chunk_alloc()被回调, 并且之后分配的地址都在该2M大的chunk内。